A Cloud Solution for Telemetry Data

Monday April 24th, 2017


Thong Bui


Roy Gvirtsman


Geoffrey Link


Zhongqiao Jin


Happiness Munedzimwe

Berkeley
UNIVERSITY OF CALIFORNIA

What is Telemetry?

- A collection of sensor reads over time
- Regular or Irregular Reads (eg., storm outages)
- Key-Value pairs: { “point”: 180, “epoch”: 1234567, “value”: 14.6}
- Examples: Temperature, Wind Direction, Alerts, GPS Coordinates

How is Telemetry used?

- Monitor and Control
- SCADA or Supervisory Control and Data Acquisition
- Alerts, Troubleshooting, Emergency Response
- Operational efficiency: KPI, Decision Support
- Predictive maintenance

KAPACITOR & SLACK

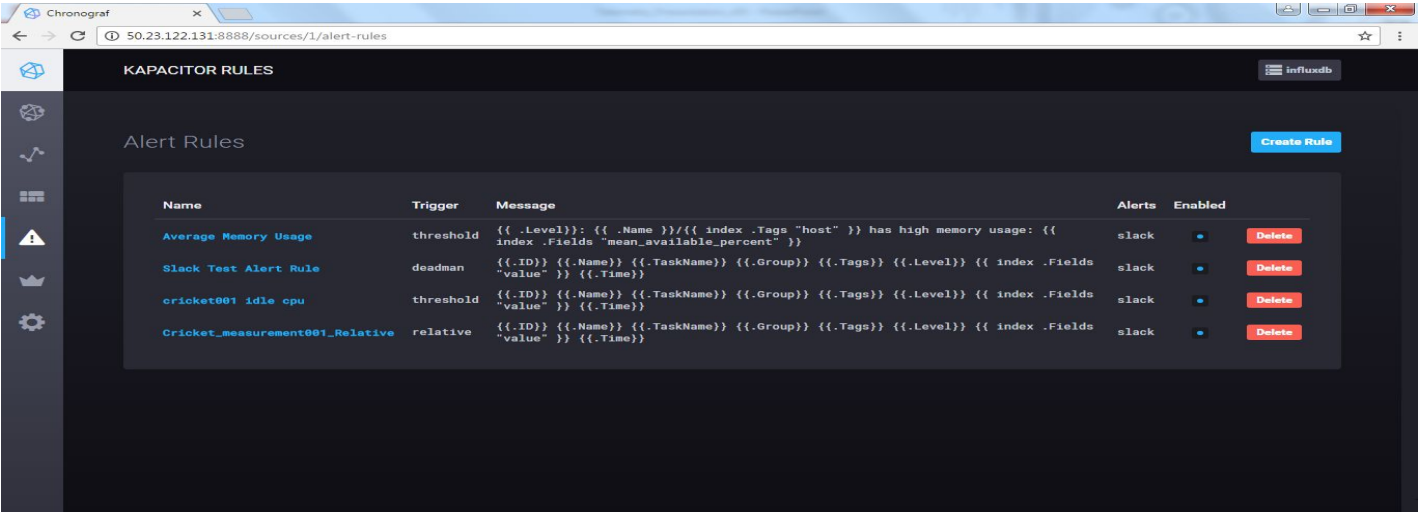

[Kapacitor Alerting Rules]

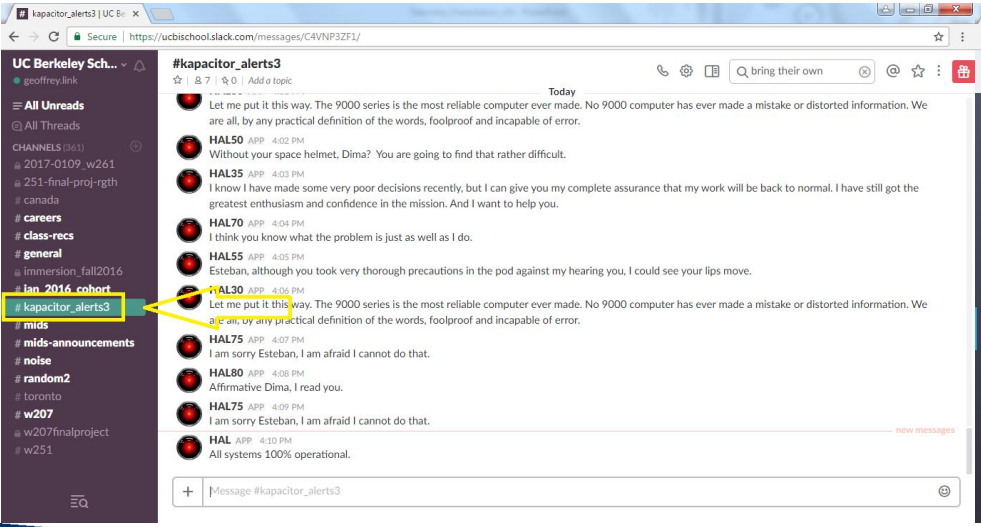

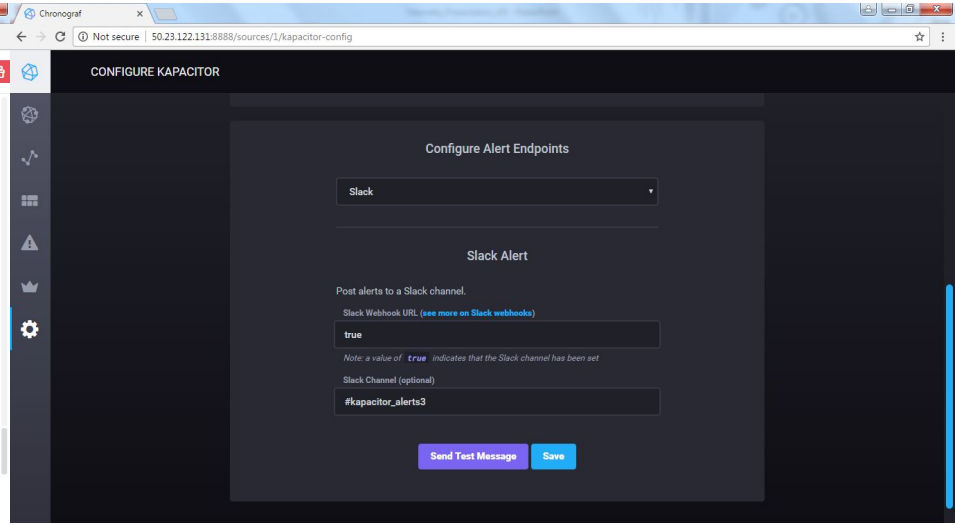

Berkeley
UNIVERSITY OF CALIFORNIA

KAPACITOR & SLACK (OpenSource Lessons Learned)

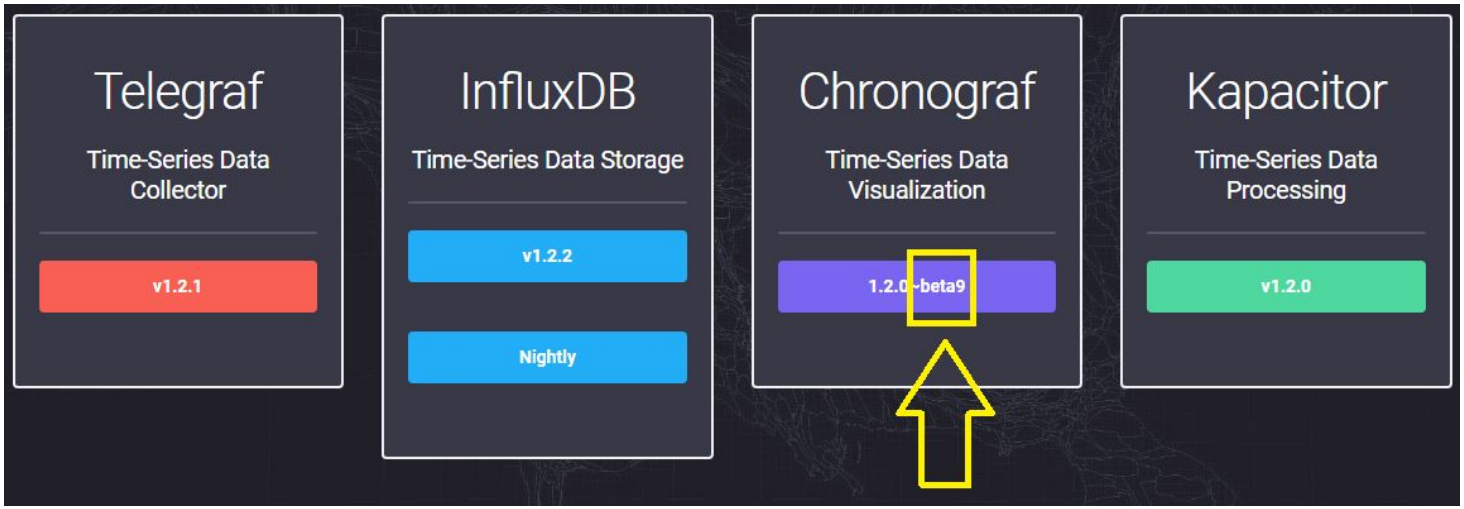

Hence, build your own SLACK alerting:

[kapacitor_alerts3 Slack Channel]

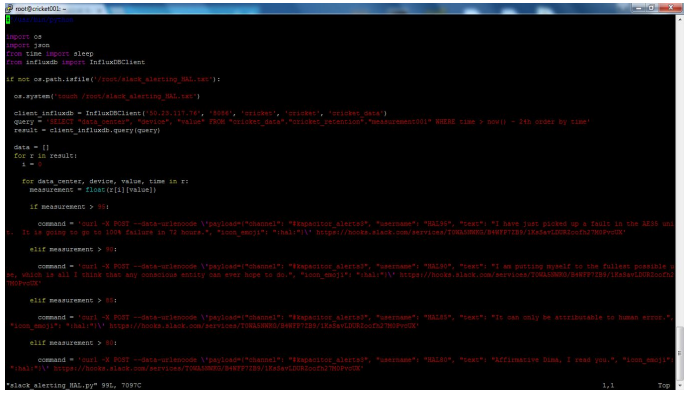

Berkeley
UNIVERSITY OF CALIFORNIA

Target Audience & Usage

Chronograf Host List

Cricket001

- Remote Operators
  - Respond to alerts
  - Troubleshoot and recover from equipment failures
- Operations Manager
  - Monitor operator efficiency
  - Quantify lost opportunity
  - Preventative maintenance
  - Handle escalated operational issues
- Market scheduler
  - Profitability
  - Respond to grid operator's instructio

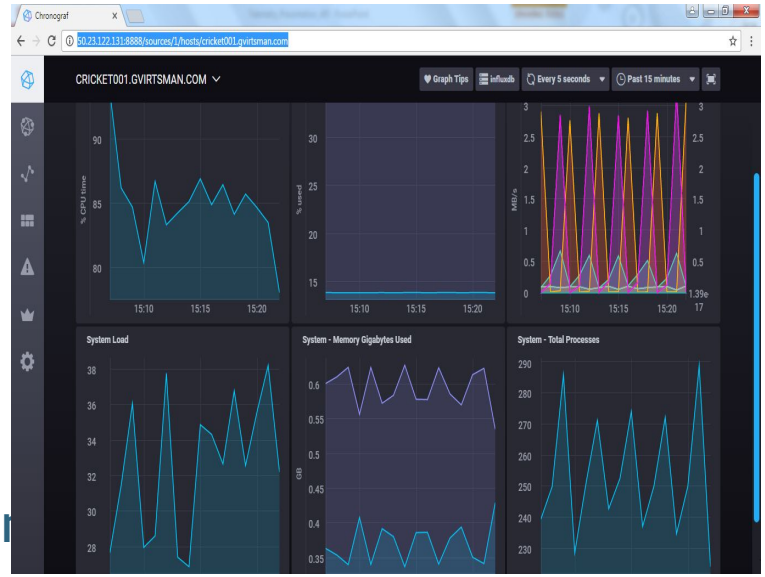

Project Objectives

- Create a reporting framework that has the ability to ingest billions of reads from tens of thousands of devices (noisy crickets)
- Easily scalable
- Fully Hosted in the Cloud (Zero Cap-Ex Platform)
- Use only open-source tools
- Each device (or cricket) will emit a basic output that can be transmitted over Radio Frequency or Wi-Fi or Cellular

Data processing requirements

- Data Generation (Python Crickets and Softlayer)
- Data Summarization (Grafana)
- Graphing and Reporting (Chronograf and D3.js)
- Retention policy (InfluxDB, 365 Days)
- Resilient and Scalable (InfluxDB, n-nodes)
- Event Subscription (Slack instead of email)
- Alerts and Notifications (Telegraf and Kapacitor)
- Predictive Maintenance (K-Means and Neural Networks)

What is a cricket? Edge Computing!

- A Cricket is a device that can transmit data over Radio Frequency or Wi-Fi or Cellular
- A common example is the highly versatile Raspberry Pi ($30).
- Raspberry Pi Sense HAT with Orientation, Pressure, Humidity, and Temperature Sensors ($30).
- In a peer-to-peer network, adding more "load" – or participants in the network – does not require any additional resources as each participant brings their own resources.

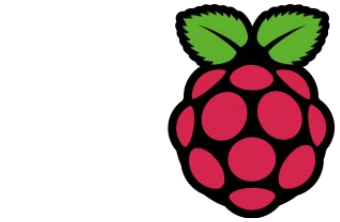

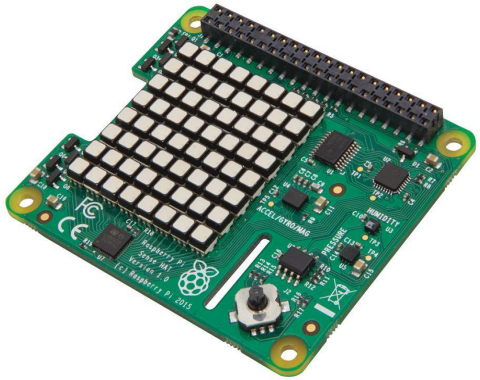

Did we buy 10,000 Raspberry Pi? No.

- We used Softlayer Cloud to simulate crickets across the globe.

Amsterdam 01	ams01	Dallas 13	dal13	San Jose 01	sjc01
Amsterdam 03	ams03	Frankfurt	fra02	San Jose 03	sjc03
Chennai	che01	Hong Kong	hkg02	Sao Paulo	sao01
Dallas 01	dal01	Houston	hou02	Seattle	sea01
Dallas 02	dal02	London	lon02	Seoul 01	seo01
Dallas 05	dal05	Melbourne	mel01	Singapore	sng01
Dallas 06	dal06	Milan	mil01	Sydney	syd01
Dallas 07	dal07	Montreal	mon01	Tokyo	tok02
Dallas 09	dal09	Oslo	osl01	Toronto	tor01
Dallas 10	dal10	Paris	par01	Washington, D.C. 01	wdc01
Dallas 12	dal12	Querétaro	mex01	Washington, D.C. 04	wdc04

```
root@cricket001: ~
root@cricket001:~# ls -l /root/cricket_message_generator.py
-rwxrwxrwx 1 root root 2638 Apr  3 00:58 /root/cricket_message_generator.py
root@cricket001:~# crontab -l
* * * * * /root/cricket_message_generator.py 50.23.117.76 cricket_001_01 sjc01 /root/cricket_001_01_data.txt > /root/cricket_001_01_message_generator.log 2>&1
* * * * * /root/cricket_message_generator.py 50.23.117.76 cricket_001_02 sjc01 /root/cricket_001_02_data.txt > /root/cricket_001_02_message_generator.log 2>&1
* * * * * /root/cricket_message_generator.py 50.23.117.76 cricket_001_03 sjc01 /root/cricket_001_03_data.txt > /root/cricket_001_03_message_generator.log 2>&1
* * * * * /root/cricket_message_generator.py 50.23.117.76 cricket_001_04 sjc01 /root/cricket_001_04_data.txt > /root/cricket_001_04_message_generator.log 2>&1
* * * * * /root/cricket_message_generator.py 50.23.117.76 cricket_001_05 sjc01 /root/cricket_001_05_data.txt > /root/cricket_001_05_message_generator.log 2>&1
* * * * * /root/cricket_message_generator.py 50.23.117.76 cricket_001_06 sjc01 /root/cricket_001_06_data.txt > /root/cricket_001_06_message_generator.log 2>&1
* * * * * /root/cricket_message_generator.py 50.23.117.76 cricket_001_07 sjc01 /root/cricket_001_07_data.txt > /root/cricket_001_07_message_generator.log 2>&1
* * * * * /root/cricket_message_generator.py 50.23.117.76 cricket_001_08 sjc01 /root/cricket_001_08_data.txt > /root/cricket_001_08_message_generator.log 2>&1
* * * * * /root/cricket_message_generator.py 50.23.117.76 cricket_001_09 sjc01 /root/cricket_001_09_data.txt > /root/cricket_001_09_message_generator.log 2>&1
* * * * * /root/cricket_message_generator.py 50.23.117.76 cricket_001_10 sjc01 /root/cricket_001_10_data.txt > /root/cricket_001_10_message_generator.log 2>&1
root@cricket001:~#
```

Berkeley
UNIVERSITY OF CALIFORNIA

Architecture Overview

C:\Windows\system32\CMD.exe

```
Microsoft Windows [Version 6.1.7601]
Copyright (c) 2009 Microsoft Corporation.  All rights reserved.

C:\Users\linkg>slcli vs list
:----------:-----------:--------------:-------------:-----------:--------:
:    id    :  hostname  :   primary_ip  :  backend_ip  : datacenter : action :
:----------:-----------:--------------:-------------:-----------:--------:
: 30427617 : chronograf : 50.23.122.131 :  10.54.136.92 :   sjc01   :   -    :
: 30167535 : cricket001 : 50.23.122.130 :  10.54.136.82 :   sjc01   :   -    :
: 29975067 :  grafana   : 169.53.133.136 : 10.122.73.100 :   sjc01   :   -    :
: 31048793 :  influx1a  : 198.11.219.238 :  10.88.154.15 :   sjc01   :   -    :
: 31048873 :  influx2a  : 50.23.100.179  : 10.54.136.113 :   sjc01   :   -    :
: 31048911 :  influx3a  : 198.11.219.236 :  10.88.154.19 :   sjc01   :   -    :
:----------:-----------:--------------:-------------:-----------:--------:

C:\Users\linkg>_
```

root@influx1a: ~

```
root@influx1a:~# influxd-ctl show
Data Nodes
==========
ID      TCP Address                     Version
2       influx1a.gvirtsman.com:8088     1.2.1-c1.2.2
5       influx3a.gvirtsman.com:8088     1.2.1-c1.2.2

Meta Nodes
==========
TCP Address                     Version
influx1a.gvirtsman.com:8091     1.2.1-c1.2.2
influx2a.gvirtsman.com:8091     1.2.1-c1.2.2
influx3a.gvirtsman.com:8091     1.2.1-c1.2.2
root@influx1a:~#
```

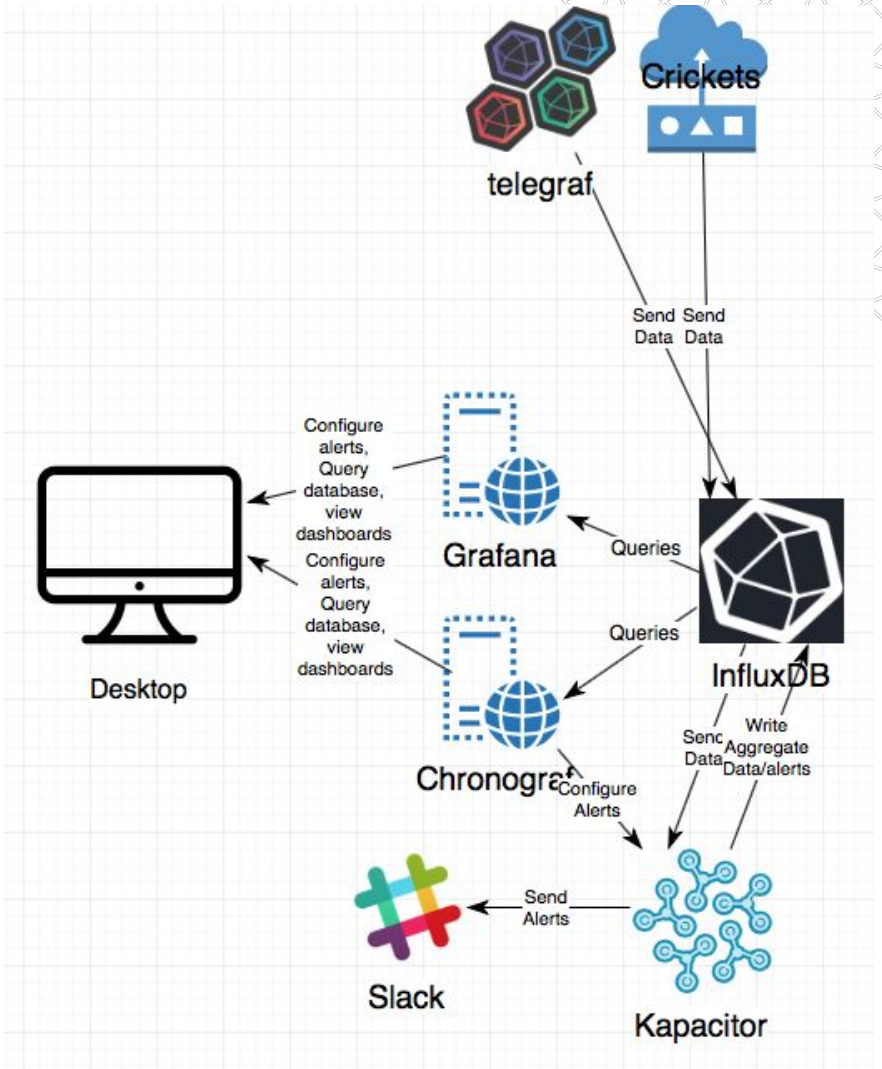

Technology Stack

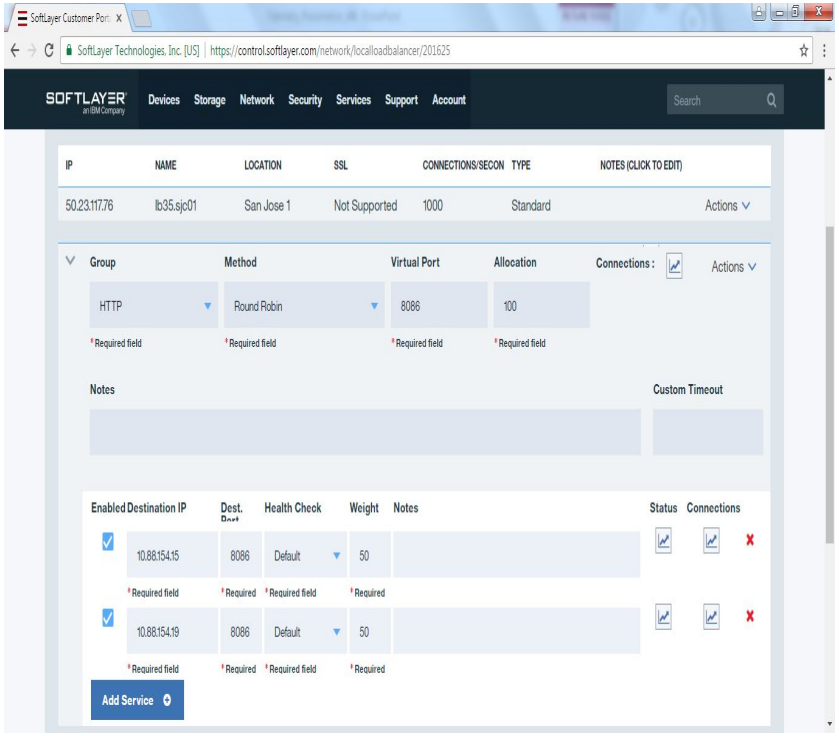

A Load Balancer is a great way to control data flow to InfluxDB as long as the crickets can store data and recover independently

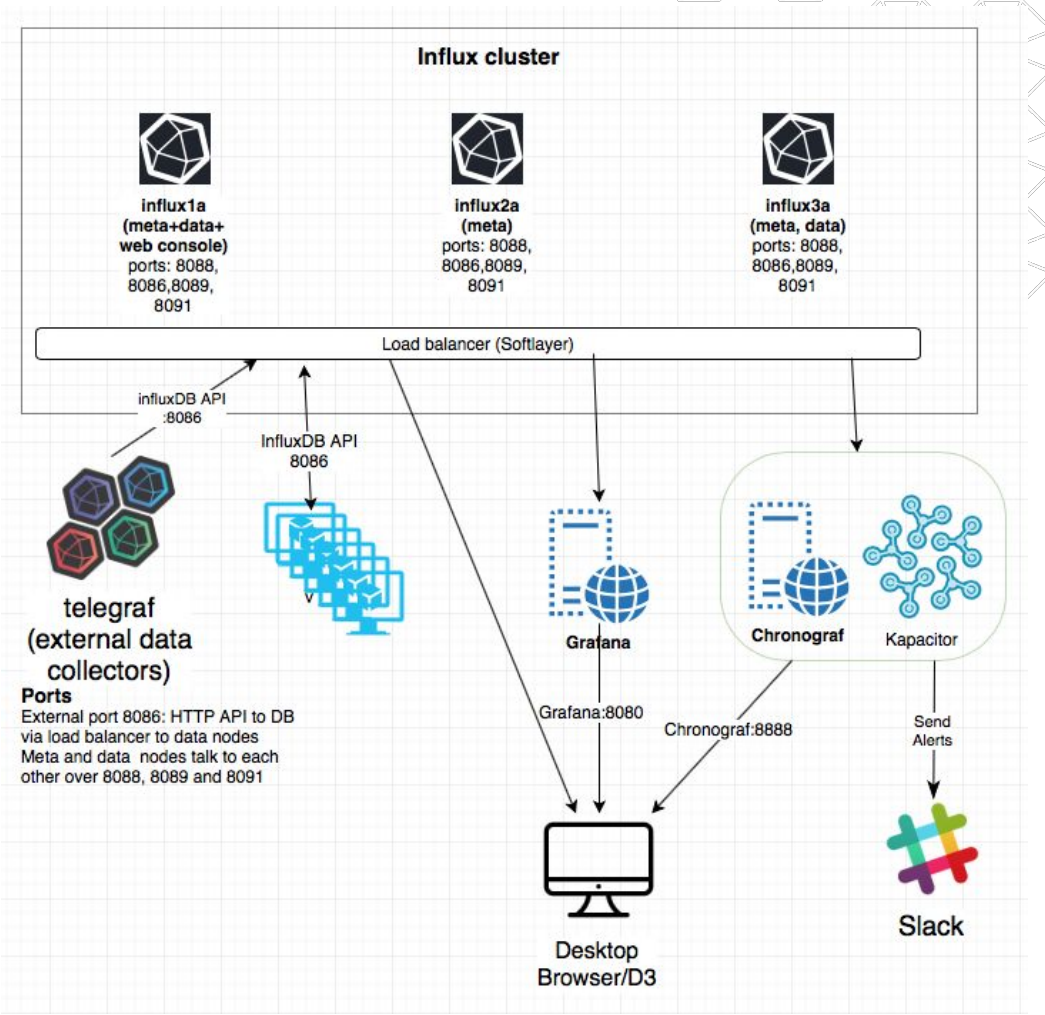

Berkeley
UNIVERSITY OF CALIFORNIA

Predictive Maintenance using K-Means

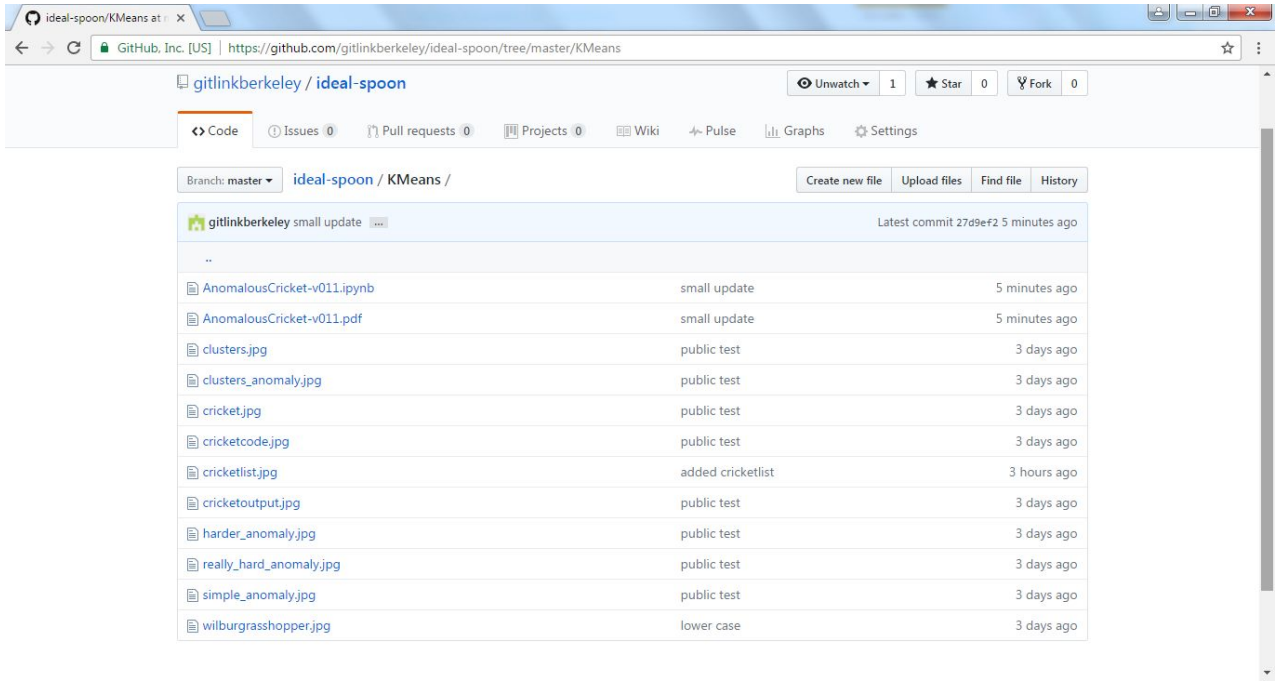

[NBViewer]

If we had more time…

- Realistic Data Construction (not randomized)
- Spark Streaming K-Means in concert with live Projected Waveforms (D3)
- Intelligent Crickets: Ex-ante Alerts followed by Self-Repair at the Edge
- Analyze Clusters of Devices and their combined Waveforms
- Bake-off Old School K-Means vs. New School Neural Networks

Berkeley
UNIVERSITY OF CALIFORNIA

InfluxDB: time series database

TICK Stack

InfluxDB

- Open source
- TICK stack (Telegraf, Influxdb, Chronograf, Kapacitor)
- Horizontally scalable (paid for product).
- 3 Meta nodes, 2+ data nodes.
- ~250,000 inserts/sec, ~25 queries/sec on a single node.
- Retention policy + # replicated copies.
- Quorum: Write: one copy, Read: choice per sessions. (biased towards heavy writes, light reads, AP)
- SQL like syntax
- Column store, compression, LSM tree (two component – memory + disk, defer index updates)

InfluxDB Query Performance

D3.js

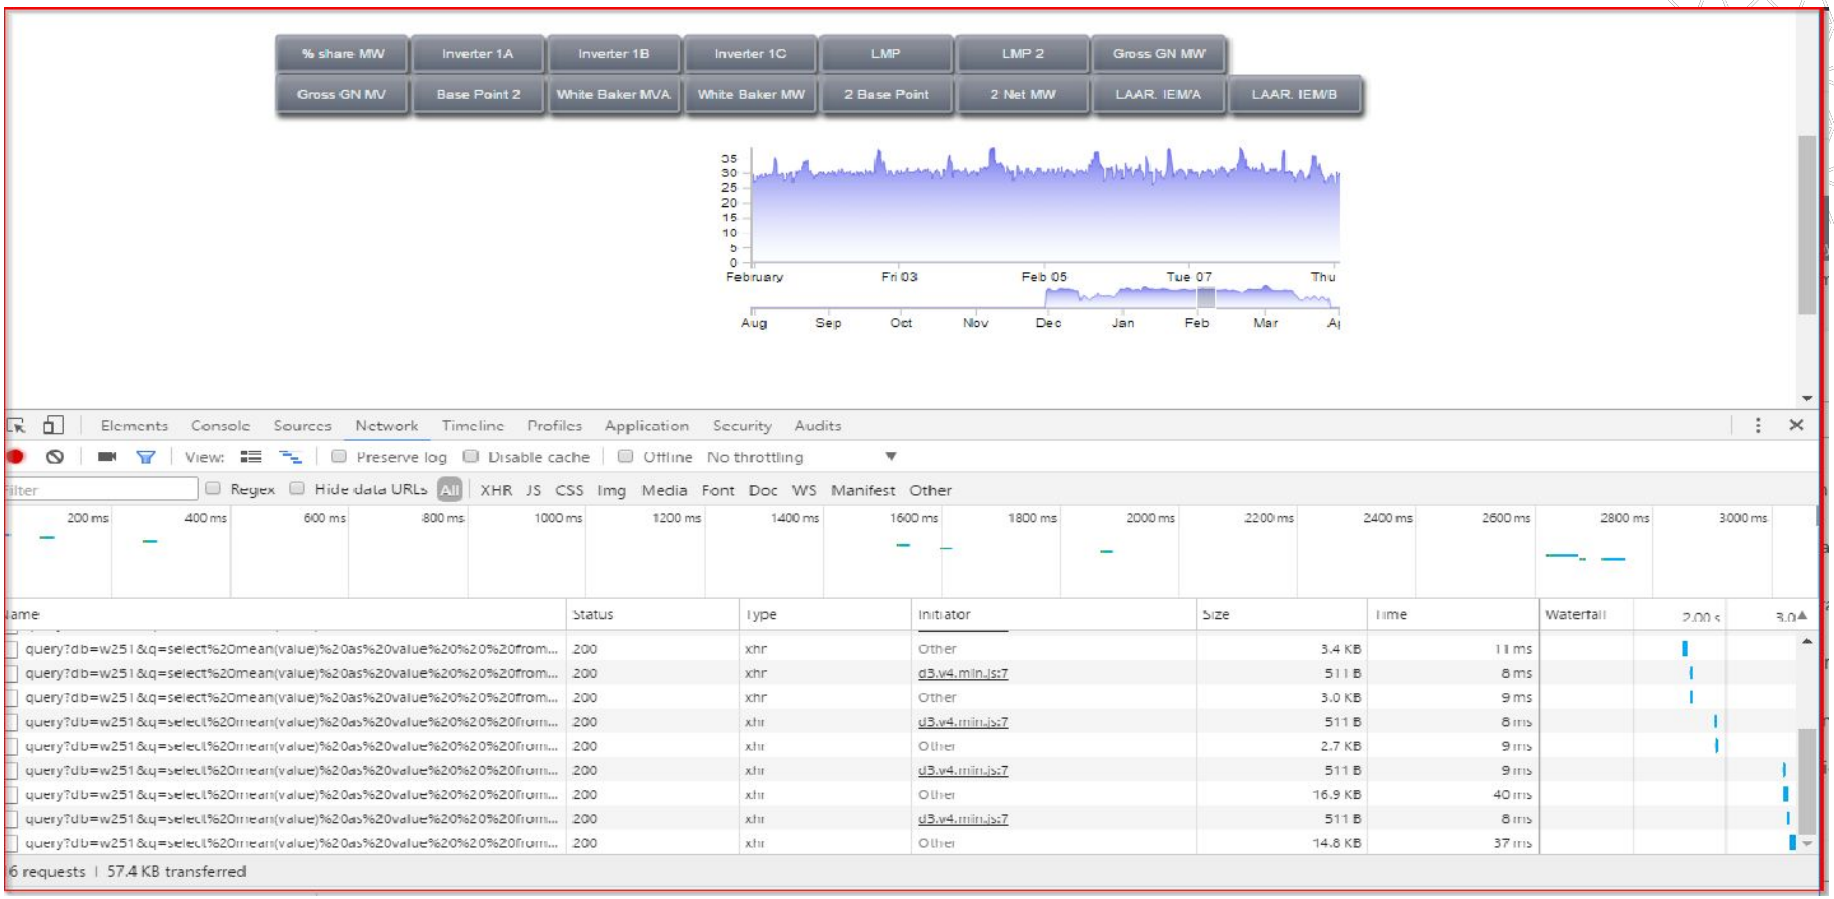

Grafana Dashboard

[Grafana Dashboard]

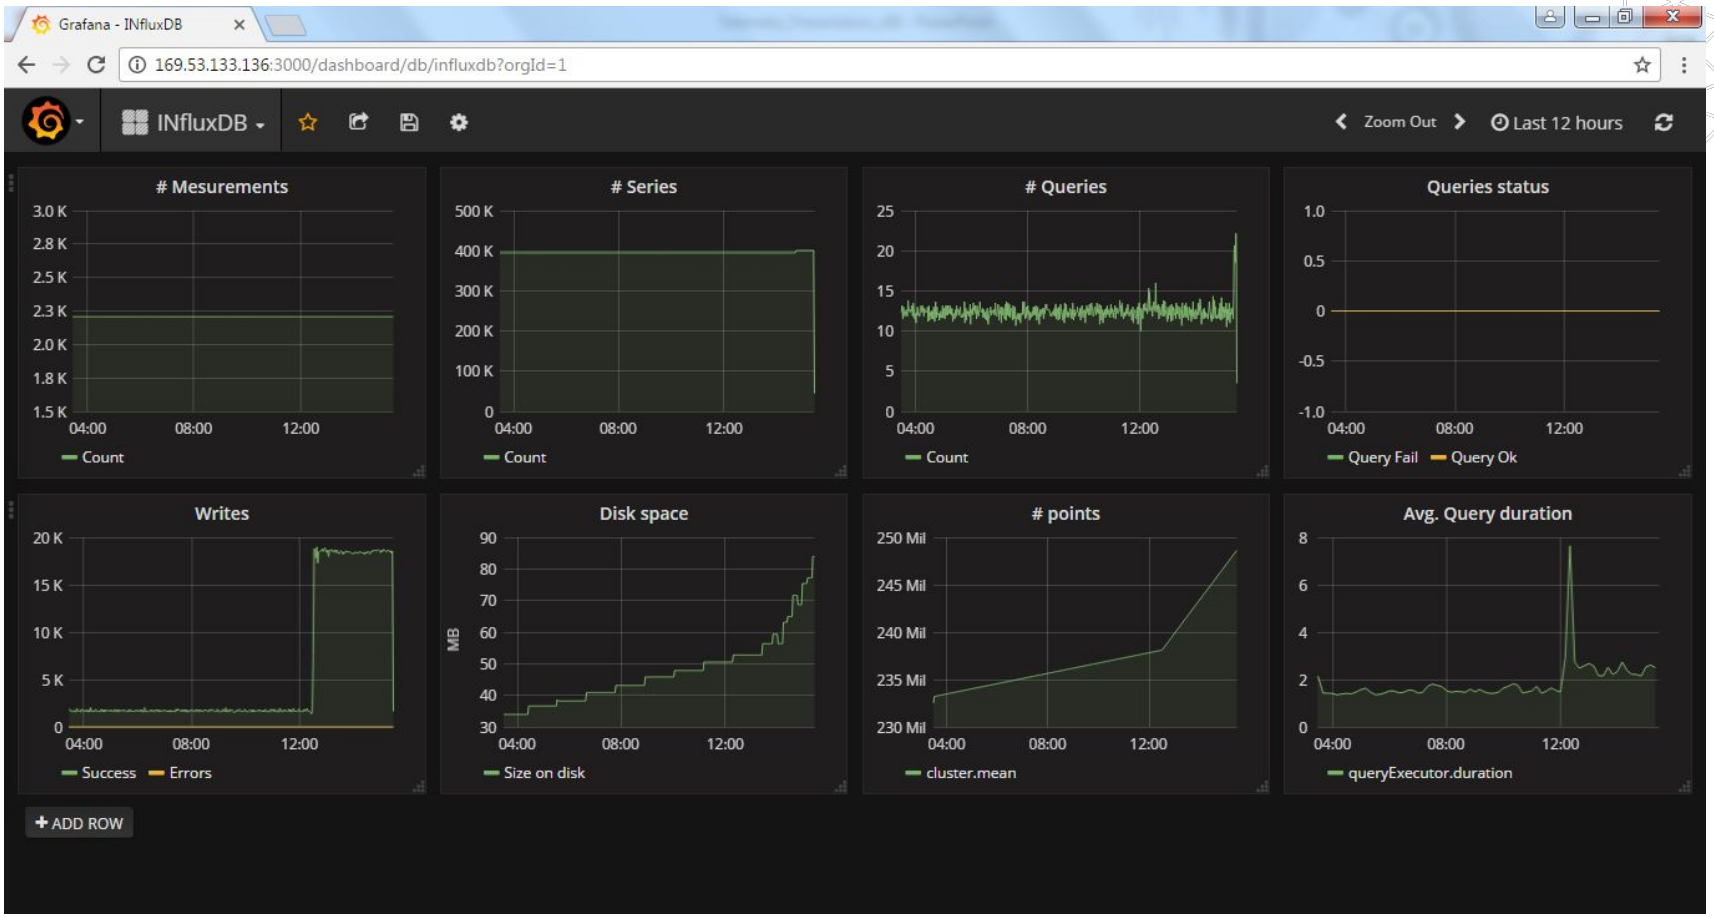

Questions?


Thong Bui


Roy Gvirtsman


Geoffrey Link


Zhongqiao Jin


Happiness Munedzimwe

Github Repository